

EXPERIMENT NO. 6 (A)

Prerequisite:

1) Download and Install Docker Desktop from <https://www.docker.com/>

Steps:

1. Check docker installation of looking at it's help page

```
quantum@machine ~$ docker -h
Flag shorthand -h has been deprecated, use --help

Usage:  docker [OPTIONS] COMMAND

A self-sufficient runtime for containers

Common Commands:
  run          Create and run a new container from an image
  exec         Execute a command in a running container
  ps           List containers
  build        Build an image from a Dockerfile
  pull         Download an image from a registry
  push         Upload an image to a registry
  images       List images
  login        Log in to a registry
  logout       Log out from a registry
  search       Search Docker Hub for images
  version      Show the Docker version information
  info         Display system-wide information

Management Commands:
  builder      Manage builds
  buildx*      Docker Buildx
  checkpoint   Manage checkpoints
  compose*     Docker Compose
  container    Manage containers
  context      Manage contexts
  dev*         Docker Dev Environments
  extension*   Manages Docker extensions
  feedback*    Provide feedback, right in your terminal!
  image        Manage images
  init*        Creates Docker-related starter files for your project
  manifest     Manage Docker image manifests and manifest lists
  network      Manage networks
  plugin       Manage plugins
  sbom*        View the packaged-based Software Bill Of Materials (SBOM) for an image
  scan*        Docker Scan
  scout*       Docker Scout
  system       Manage Docker
  trust        Manage trust on Docker images
  volume       Manage volumes
```

```
quantum@machine ~ docker -v
Docker version 27.1.2, build d01f264
```

2. Create a new folder named **"Terraform Scripts"**, inside it create a new folder **docker** and create a new file named **docker.tf** with following contents inside it.

```
GNU nano 7.2 docker.tf
terraform {
  required_providers {
    docker = {
      source = "kreuzwerker/docker"
      version = "2.21.0"
    }
  }
}

provider "docker" {
  host = "unix:///var/run/docker.sock"
}

# Pulls the image
resource "docker_image" "ubuntu" {
  name = "ubuntu:latest"
}

# Create a container
resource "docker_container" "foo" {
  image = docker_image.ubuntu.image_id
  name = "foo"
  command = ["/bin/bash", "-c", "while true; do sleep 3600; done"]
}
```

3. To see list of providers for current configuration file use command **terraform providers**

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker terraform providers

Providers required by configuration:
└─ provider[registry.terraform.io/kreuzwerker/docker] 2.21.0
```

4. Execute command **terraform init** in the current directory

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker terraform init
Initializing the backend...
Initializing provider plugins...
- Finding kreuzwerker/docker versions matching "2.21.0"...
- Installing kreuzwerker/docker v2.21.0...
- Installed kreuzwerker/docker v2.21.0 (self-signed, key ID BD080C4571C6104C)
Partner and community providers are signed by their developers.
If you'd like to know more about provider signing, you can read about it here:
https://www.terraform.io/docs/cli/plugins/signing.html
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

5. Use command **terraform validate** to check for validation and syntax errors of config file

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker terraform validate
Success! The configuration is valid.
```

5. Execute command **terraform plan** to see the changes that will be made

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker$ sudo terraform plan
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# docker_container.foo will be created
+ resource "docker_container" "foo" {
  + attach      = false
  + bridge      = (known after apply)
  + command     = (known after apply)
  + container_logs = (known after apply)
  + entrypoint  = (known after apply)
  + env         = (known after apply)
  + exit_code   = (known after apply)
  + gateway     = (known after apply)
  + hostname    = (known after apply)
  + id          = (known after apply)
  + image       = (known after apply)
  + init        = (known after apply)
  + ip_address  = (known after apply)
  + ip_prefix_length = (known after apply)
  + ipc_mode    = (known after apply)
  + log_driver  = (known after apply)
  + logs        = false
  + must_run    = true
  + name        = "foo"
  + network_data = (known after apply)
  + read_only   = false
  + remove_volumes = true
  + restart     = "no"
  + rm          = false
  + runtime     = (known after apply)
  + security_opts = (known after apply)
  + shm_size    = (known after apply)
  + start       = true
  + stdin_open  = false
  + stop_signal  = (known after apply)
  + stop_timeout = (known after apply)
  + tty         = false
  + healthcheck (known after apply)
}
```

```

+ labels (known after apply)
}

# docker_image.ubuntu will be created
+ resource "docker_image" "ubuntu" {
  + id          = (known after apply)
  + image_id    = (known after apply)
  + latest      = (known after apply)
  + name        = "ubuntu:latest"
  + output      = (known after apply)
  + repo_digest = (known after apply)
}

Plan: 2 to add, 0 to change, 0 to destroy.
```

6. Check list of running docker containers using command **docker ps**

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

As we can see there are no active containers running

7. Execute **terraform apply** to apply configuration, which will automatically create and run the Ubuntu Linux container based on our configuration.

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker sudo terraform apply
docker_image.ubuntu: Refreshing state... [id=sha256:17c0145030df106e60e5d99149d69810db23b869ff0d3c9d23627a5a7bbb6b3ubuntu:latest]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# docker_container.foo will be created
+ resource "docker_container" "foo" {
  + attach      = false
  + bridge      = (known after apply)
  + command     = [
    + "/bin/bash",
    + "-c",
    + "while true; do sleep 3600; done",
  ]
  + container_logs = (known after apply)
  + entrypoint     = (known after apply)
  + env           = (known after apply)
  + exit_code      = (known after apply)
  + gateway        = (known after apply)
  + hostname       = (known after apply)
  + id             = (known after apply)
  + image          = "sha256:17c0145030df106e60e5d99149d69810db23b869ff0d3c9d236279a5a7bbb6b3"
  + init          = (known after apply)
  + ip_address     = (known after apply)
  + ip_prefix_length = (known after apply)
  + ipc_mode       = (known after apply)
}

Plan: 1 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

Enter a value: yes

docker_container.foo: Creating...
docker_container.foo: Creation complete after 1s [id=d912cf0a2579f9b8c958d2d33426ed11e8251a553f85a0c08022d19ddf9eeecd]

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

8. After executing **terraform apply**

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker$ sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
d912cf0a2579	17c0145030df	"/bin/bash -c 'while..."	2 minutes ago	Up About a minute		foo

After execution, a new docker image and container will be created

9. Execute **terraform destroy** to delete the configuration, which will automatically delete the Container.

```
# docker_image.ubuntu will be destroyed
- resource "docker_image" "ubuntu" {
  - id          = "sha256:35a88802559dd2077e584394471ddaa1a2c5bfd16893b829ea57619301eb3908ubuntu:latest"
  -> null
  - image_id    = "sha256:35a88802559dd2077e584394471ddaa1a2c5bfd16893b829ea57619301eb3908" -> null
  - latest      = "sha256:35a88802559dd2077e584394471ddaa1a2c5bfd16893b829ea57619301eb3908" -> null
  - name        = "ubuntu:latest" -> null
  - repo_digest = "ubuntu@sha256:2e863c44b718727c860746568e1d54afd13b2fa71b160f5cd9058fc436217b30" -> null
}

Plan: 0 to add, 0 to change, 2 to destroy.

Do you really want to destroy all resources?
Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

docker_container.foo: Destroying... [id=8329fd298f0ccf122a391d74108a2809c40b5fbc383ffc31ee4fb6aeb6f0e3c7]
docker_container.foo: Destruction complete after 0s
docker_image.ubuntu: Destroying... [id=sha256:35a88802559dd2077e584394471ddaa1a2c5bfd16893b829ea57619301eb3908ubuntu:latest]
docker_image.ubuntu: Destruction complete after 0s

Destroy complete! Resources: 2 destroyed.
```

10. Check whether docker image and container is removed or not

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker$ sudo docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
alpine	latest	05455a08881e	6 months ago	7.38MB
busybox	latest	3f57d9401f8d	7 months ago	4.26MB
sxcurity/gau	latest	5a5fc3cf7aa4	9 months ago	23.3MB
hello-world	latest	d2c94e258dcb	15 months ago	13.3kB

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/docker$ sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

As we can see both containers and images are removed

EXPERIMENT NO. 6 (B)

Prerequisite:

- 1) Any text editor to write and save scripts
- 2) Must have an AWS Access Key ID and Secret Access Key

Step 1: Write a Terraform Script in Atom for creating S3 Bucket on Amazon AWS

Ensure your bucket name is globally unique

Step 2: Get secret key and access key from AWS account

1. Go to IAM dashboard

The screenshot displays the AWS Identity and Access Management (IAM) console. On the left, a sidebar contains the 'Identity and Access Management (IAM)' header, a search bar labeled 'Search IAM', and a navigation menu with 'Dashboard' (highlighted in blue) and 'Access management' (expanded to show 'User groups', 'Users', 'Roles', 'Policies', 'Identity providers', and 'Account settings'). The main content area is titled 'IAM Dashboard' and includes a breadcrumb 'IAM > Dashboard'. Below the title, there are two primary sections: 'Security recommendations' and 'IAM resources'. The 'Security recommendations' section features a green status indicator with the number '0' and a refresh button. It lists two recommendations, both marked with green checkmarks: 'Root user has MFA' (with a sub-note: 'Having multi-factor authentication (MFA) for the root user improves security for this account.') and 'Root user has no active access keys' (with a sub-note: 'Using access keys attached to an IAM user instead of the root user improves security.'). The 'IAM resources' section is partially visible at the bottom, showing the title and a refresh button.

oll Go to Users and select anNo:59 existingiv user if it exists or create a new one. I'm selecting existing one as i had already created before

Users (1) [Info](#)

[Refresh](#) [Delete](#) [Create user](#)

An IAM user is an identity with long-term credentials that is used to interact with AWS in an account.

< 1 >

<input type="checkbox"/>	User name	Path	Group:	Last activity	MFA	Password age	Console last sign-in	Access key ID	Active key age
<input type="checkbox"/>	aloky	/	0	-	-	-	-	-	-

3. On the user dashboard, click on **create access key**

aloky [Info](#)

[Delete](#)

Summary

ARN arn:aws:iam::058264551418:user/aloky	Console access Disabled	Access key 1 Create access key
Created August 12, 2024, 20:39 (UTC+05:30)	Last console sign-in -	

4. Select **Use case** as **Local Code** then Create access key

Use case

☐ Command Line Interface (CLI)
You plan to use this access key to enable the AWS CLI to access your AWS account.

☒ Local code
You plan to use this access key to enable application code in a local development environment to access your AWS account.

☐ Application running on an AWS compute service
You plan to use this access key to enable application code running on an AWS compute service like Amazon EC2, Amazon ECS, or AWS Lambda to access your AWS account.

☐ Third-party service
You plan to use this access key to enable access for a third-party application or service that monitors or manages your AWS resources.

☐ Application running outside AWS
You plan to use this access key to authenticate workloads running in your data center or other infrastructure outside of AWS that needs to access your AWS resources.

☐ Other
Your use case is not listed here.

5. After all above steps, access key is generated

is key

Retrieve access keys [Info](#)

Access key

If you lose or forget your secret access key, you cannot retrieve it. Instead, create a new access key and make the old key inactive.

Access key	Secret access key
AKIAQ3EGWP75IK42V5N3	VekODARrrlaVDCPbsTqLbHdOIQ+kc5HGfSeC9Cu7 Hide

Access key best practices

- Never store your access key in plain text, in a code repository, or in code.
- Disable or delete access key when no longer needed.
- Enable least-privilege permissions.
- Rotate access keys regularly.

For more details about managing access keys, see the [best practices for managing AWS access keys](#).

[Download .csv file](#)
[Done](#)

- Ensure newly created user has necessary permissions to edit and create S3 buckets to do so, click on **Add permissions**

Permissions policies (1) [Refresh](#) [Remove](#) [Add permissions](#)

Permissions are defined by policies attached to the user directly or through groups.

- Select permission option as **Attach policies directly** and select **AmazonS3FullAccess** as permission policy from the list

Add permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

Permissions options

☐ **Add user to group**
Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.

☐ **Copy permissions**
Copy all group memberships, attached managed policies, inline policies, and any existing permissions boundaries from an existing user.

☒ **Attach policies directly**
Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.

Permissions policies

 (1/1228) [Refresh](#)

12 matches

<input type="checkbox"/>	Policy name	Type	Attached entities
<input type="checkbox"/>	AmazonDMSRedshiftS3Role	AWS managed	0
<input checked="" type="checkbox"/>	AmazonS3FullAccess	AWS managed	0
<input type="checkbox"/>	AmazonS3ObjectLambdaExecutionRolePolicy	AWS managed	0

Step 3: Create a new provider.tf file and write the following contents into it.

```
provider "aws" {  
  access_key = "AKIAQ3EGWP75IK42VSN3"  
  secret_key = "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"  
  region     = "ap-south-1"  
}
```

Save both the files in same directory TerraformScripts/S3

Step 4: Open terminal and go to TerraformScripts/S3 directory where our .tf files are stored

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/S3  
ls  
provider.tf s3.tf
```

Step 5: Execute **terraform init** to initialize the resources

```
quantum@machine ~/Downloads/advdevops/TerraformScripts/S3 terraform init  
Initializing the backend...  
Initializing provider plugins...  
- Finding latest version of hashicorp/aws...  
- Installing hashicorp/aws v5.63.1...  
- Installed hashicorp/aws v5.63.1 (signed by HashiCorp)  
Terraform has created a lock file .terraform.lock.hcl to record the provider  
selections it made above. Include this file in your version control repository  
so that Terraform can guarantee to make the same selections by default when  
you run "terraform init" in the future.  
  
Terraform has been successfully initialized!  
  
You may now begin working with Terraform. Try running "terraform plan" to see  
any changes that are required for your infrastructure. All Terraform commands  
should now work.  
  
If you ever set or change modules or backend configuration for Terraform,  
rerun this command to reinitialize your working directory. If you forget, other  
commands will detect it and remind you to do so if necessary.
```

Step 6: Execute Terraform plan to see the available resources

```

quantum@machine ~/Downloads/advdevops/TerraformScripts/S3: sudo terraform plan
[sudo] password for quantum:

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_s3_bucket.kajal will be created
+ resource "aws_s3_bucket" "kajal" {
  + acceleration_status      = (known after apply)
  + acl                      = "public-read"
  + arn                      = (known after apply)
  + bucket                  = "my-bj-terraform-test-bucket"
  + bucket_domain_name      = (known after apply)
  + bucket_prefix           = (known after apply)
  + bucket_regional_domain_name = (known after apply)
  + force_destroy           = false
  + hosted_zone_id          = (known after apply)
  + id                      = (known after apply)
  + object_lock_enabled      = (known after apply)
  + policy                  = (known after apply)
  + region                  = (known after apply)
  + request_payer            = (known after apply)
  + tags                    = {
    + "Environment" = "Dev"
    + "Name"        = "My bucket"
  }
  + tags_all              = {
    + "Environment" = "Dev"
    + "Name"        = "My bucket"
  }
  + website_domain         = (known after apply)
  + website_endpoint       = (known after apply)

  + cors_rule (known after apply)

  + grant (known after apply)

  + lifecycle_rule (known after apply)

```

```

+ logging (known after apply)

+ object_lock_configuration (known after apply)

+ replication_configuration (known after apply)

+ server_side_encryption_configuration (known after apply)

+ versioning (known after apply)

+ website (known after apply)
}

Plan: 1 to add, 0 to change, 0 to destroy.

Warning: Argument is deprecated

with aws_s3_bucket.kajal,
on s3.tf line 3, in resource "aws_s3_bucket" "kajal":
3:   acl = "public-read"

Use the aws_s3_bucket_acl resource instead

(and one more similar warning elsewhere)

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

```

Step 7: Execute **terraform apply** to apply the configuration, which will automatically create an S3 bucket based on our configuration.

```
quantum@machine ~/Downloads/advdevops/terraformScripts/S3 > sudo terraform apply
```

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

```
# aws_s3_bucket.aloky will be created
+ resource "aws_s3_bucket" "aloky" {
  + acceleration_status = (known after apply)
  + acl                  = (known after apply)
  + arn                  = (known after apply)
  + bucket               = "my-bj-terraform-test-bucket-aloky"
  + bucket_domain_name  = (known after apply)
  + bucket_prefix        = (known after apply)
  + bucket_regional_domain_name = (known after apply)
  + force_destroy        = false
  + hosted_zone_id       = (known after apply)
  + id                   = (known after apply)
  + object_lock_enabled  = (known after apply)
  + policy               = (known after apply)
  + region               = (known after apply)
  + request_payer        = (known after apply)
  + tags                 = {
    + "Environment" = "Dev"
    + "Name"        = "My bucket"
  }
  + tags_all = {
    + "Environment" = "Dev"
    + "Name"        = "My bucket"
  }
  + website_domain      = (known after apply)
  + website_endpoint    = (known after apply)

  + cors_rule (known after apply)

  + grant (known after apply)
```

```
+ grant (known after apply)

+ lifecycle_rule (known after apply)

+ logging (known after apply)

+ object_lock_configuration (known after apply)

+ replication_configuration (known after apply)

+ server_side_encryption_configuration (known after apply)

+ versioning (known after apply)

+ website (known after apply)
}
```

Plan: 1 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

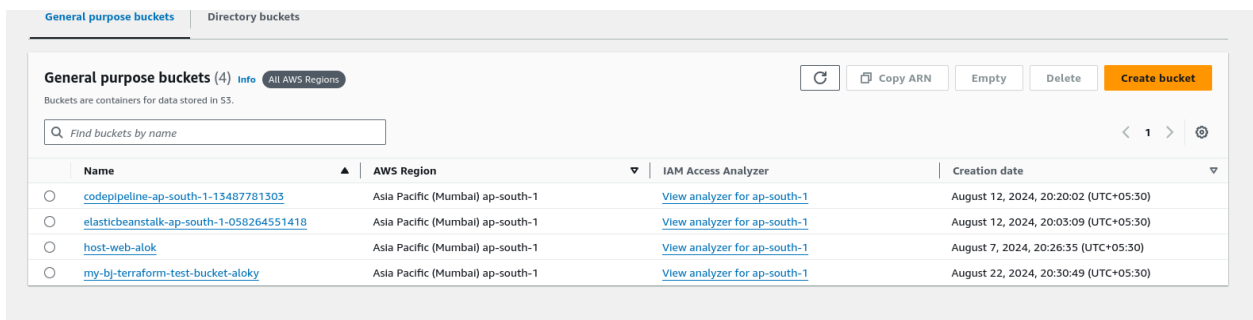
Enter a value: yes

aws_s3_bucket.aloky: Creating...

aws_s3_bucket.aloky: Creation complete after 4s [id=my-bj-terraform-test-bucket-aloky]

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

Step 8: Go to **S3 Buckets** dashboard to see newly created S3 bucket using terraform



my-bj-terraform-test-bucket-aloky Info

- Objects
- Properties
- Permissions
- Metrics
- Management
- Access Points

Objects (0) Info



Copy S3 URI

Copy URL

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For more information, see [Amazon S3 inventory](#).

Find objects by prefix

	Name	▲	Type	▼	Last modified
--	------	---	------	---	---------------

No objects